

Concept Drift, Continuous ML, MLOps & CCAM: A Drift-aware Platform

Angelis Georgios, Kyriakopoulos Georgios, Tzelepis Serafeim

School of Electrical and Computer Engineering, National Technical University of Athens

28 July 2025

Agenda

- 1) Thesis Scope
- 2) Problem Statement
- 3) Software Design & Tools
- 4) Data Acquisition
- 5) Implementation
- 6) Platform Overview
- 7) Implementation Challenges

Thesis Scope

- ▶ Build a **drift-aware platform** for mobility models on Athens center.
- ▶ Generate **synthetic traffic data** from SUMO (Simulation of Urban MObility).
- ▶ **Three regression tasks**: ETA, Fuel Consumption, Number of Stops.
- ▶ Show end-to-end: **stable** → **drift** → **detect** → **retrain** → **mitigate**.
- ▶ Deliverables: dataset generation pipeline, tuned models for each task, platform (backend, frontend, drift service, model services).

Concept Drift & Why It Matters

- ▶ **What it is:** When the real world changes, the link between inputs and the target also changes over time (non-stationarity).
- ▶ **Why it happens:** Seasons and trends, one-off events, policy or product changes, and even sensor or software updates.
- ▶ **Why it matters:** If we don't watch for it, model accuracy quietly drops, dashboards stop matching reality, and decisions become costly or unsafe.
- ▶ **What to do:** Treat ML as a live system—**monitor** data and outcomes, **detect** changes early, and **adapt**.

Our ML Tasks

- ▶ **Estimated Time of Arrival (ETA)**: essential for accurate route planning and improving user satisfaction.
- ▶ **Fuel Consumption**: key to reducing operating costs and CO₂ emissions via eco-routing.
- ▶ **Number of Stops**: influences perceived trip quality and helps optimize overall traffic flow.

Software Design & Tools

- ▶ **Synthetic Data:** SUMO - Python.
- ▶ **Models:** XGBoost/LightGBM/CatBoost, scikit-learn, numpy, pandas - Python.
- ▶ **Backend:** FastAPI - Python.
- ▶ **Drift:** River (ADWIN, KSWIN, PageHinkley) - Python.
- ▶ **UI:** Dash/Plotly - Python.
- ▶ **Data:** Parquet/CSV files.
- ▶ **Orchestration:** Docker.

Athens Network Configurations & Traffic Patterns

- ▶ **Base network:** SUMO network built from OSM data over central Athens.
- ▶ **Rain network:** same map area, but with friction reduced in all lanes ($1.0 \rightarrow 0.4$) to simulate wet conditions.
- ▶ **Traffic generation periods:** per-hour volumes with morning and afternoon peaks, and lower mid-day activity.
- ▶ **Realistic variability:** added random noise to hourly volumes and used a distinct random seed for each 10h day simulation.

Dataset Structure

- ▶ **Three CSV files** (each $\sim 55\text{--}60$ k trips over a 10 hour simulation):
 - ▶ `train-fcd.csv` (base network)
 - ▶ `test-fcd.csv` (base network)
 - ▶ `rain-fcd.csv` (rain-modified network)
- ▶ **Columns (per row)**: `timestep`, `vehicle_id`, `x`, `y`, `speed`, `lane`, `odometer`, `fuel`, `waiting_time`

Dataset Generation Pipeline

- ▶ **osmGet**: extract OSM data within a bounding box over central Athens
- ▶ **osmBuild**: convert OSM into a SUMO-compatible network file
- ▶ **randomTrips**: generate vehicle trips using per-hour traffic volumes and a different random seed
- ▶ **sumo**: run 10-hour simulations, export Floating Car Data (FCD) and emission logs

Data Preprocessing & Feature Engineering

- ▶ **Trip extraction:** group raw FCD by vehicle into one row per trip (origin-destination)
- ▶ **Different features** added per model:
 - ▶ *Temporal:* hour_bin
 - ▶ *Spatial:* source/destination clusters (x,y)
 - ▶ *Route:* number of edges traversed, distance measures (Euclidean, Manhattan, Haversine)
- ▶ **Result:** tabular dataset with inputs + ground-truth targets (ETA, fuel, stops)


Model Training & Evaluation

- ▶ **Evaluation metrics:**
 - ▶ *Primary:* MAE
 - ▶ *Secondary:* MAPE, R^2 , RMSE, MSE
- ▶ **Hyperparameter tuning:** Optuna / grid search over `n_estimators`, `max_depth`, `learning_rate`, `subsample`, ...
- ▶ **Outcome:** XGBoost delivered the best trade-off of accuracy and training speed for all three tasks

High-level Architecture

- ▶ **Frontend (Dash/Plotly)**: Admin tab (live MAE charts, notifications, run controls); User tab (Athens map, O/D query).
- ▶ **Orchestrator (FastAPI)**: simulation clock, batch prediction, error calc, drift ensemble, per-model state, UI feed, controls.
- ▶ **Model Service (FastAPI)**: per-task pipelines (preprocessing + booster), /predict, /retrain, /load.
- ▶ **Filesystem Store**: Parquet datasets (train/test/train); filesystem model registry (models/{task}/vN/).
- ▶ **Docker Compose**: orchestration with multiple containers.

Components Diagram



How It Works (at a glance)

1. **Timelapse clock:** compress 20 h of trips (test then rain) into a few minutes.
2. **Per tick:** Orchestrator pulls trips whose `sim_time` falls in the current window.
3. **Batch predict:** send features to Model Service; receive predictions for ETA/Fuel/Stops.
4. **Immediate ground truth:** join with labels; compute per-trip absolute errors.
5. **Live metrics:** update rolling MAE per task and push to the UI feed.
6. **Drift monitoring:** feed a smoothed error stream into ADWIN, Page-Hinkley, KSWIN, and an SPC rule; raise drift on majority vote.
7. **Mitigation:** after drift is confirmed, collect N more “rain” trips, trigger retraining, then swap to the new model version.
8. **Repeat per model:** ETA, Fuel, Stops progress through the same state machine independently.

Admin Tab: Sped-up Simulation & Drift Story

- ▶ **Three live charts** (ETA/Fuel/Stops): rolling MAE; drift state (Stable/Confirmed/Collecting/Retraining/Swapped).
- ▶ **Notifications**: day transition at 10 h (*rain starts*), drift detected ($\text{votes} \geq 3$), retrain started/completed, model swapped to vN .
- ▶ **What you see**: errors rise after rain; drift fires; we keep predicting with old model while “collecting” rain data; retrain runs; swap reduces error (not all the way to baseline).
- ▶ **Controls**: Start/Pause/Resume/Restart.


Admin Tab: Sped-up Simulation & Drift Story



User Tab: Predictions at the Current State

- ▶ **Map of central Athens:** user picks **source** and **destination**.
- ▶ **Uses the same simulation time:** features derive from current `sim_time` (e.g., `hour_bin`, `distance`, `spatial bins`).
- ▶ **Calls active models:** returns ETA, Fuel, Stops from whichever version is currently loaded (pre- or post-swap).
- ▶ **Consistent context:** shows badges like “Rain scenario” and current drift state so predictions match what the Admin tab shows.
- ▶ **Demo behavior:** *pauses* the timelapse while in the User tab, then *resumes* when you go back to Admin.

User Tab: Predictions at the Current State



Implementation Challenges

- ▶ **Road closures drift scenario:** Closing lanes/edges to mimic events/roadworks in SUMO led to unstable routing and distorted traffic patterns. We therefore used the rain/friction shift as a controlled, repeatable drift.
- ▶ **Parking availability task:** The current SUMO setup lacks realistic parking supply/occupancy and consistent stop events, so the task could not be integrated or evaluated alongside the trip-level models.
- ▶ **Time-series traffic speed model:** A lagged-input traffic speed model adapted quickly to the rain shift and showed little error increase, making drift hard to showcase; we replaced it with the “number of stops” task which reacts clearly.

Questions

Thank you!